

Convolutional neural networks

E. Decencière

MINES ParisTech
PSL Research University
Center for Mathematical Morphology



1 / 59

Contents

- 1 Introduction
- 2 Application of fully-connected networks to image classification
- 3 From fully-connected layers to convolutional layers
- 4 Building convolutional networks
- 5 Conclusion

2 / 59

Contents

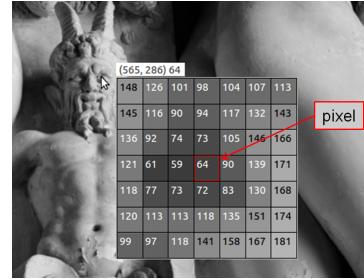
- 1 Introduction
- 2 Application of fully-connected networks to image classification
- 3 From fully-connected layers to convolutional layers
- 4 Building convolutional networks
- 5 Conclusion

3 / 59

A picture is worth a thousand words

Definition

- Classically, an image is a matrix of values belonging to $[0, \dots, 255]$ (grey level images) or to $[0, \dots, 255]^3$ (color images).
- More generally, an image is a q -dimensional array of values belonging to R^d .



Grey level values around the left eye of the faun

4 / 59

Extracting semantic information from an image



- Where is the phone? (localization task)
- How many mugs are there? (quantification task)
- Is there a window in the room?
- At what time of the day was the photograph taken?

Designing computer vision systems that are able to extract semantic information from an image is a difficult task.

5 / 59

The role of annotated image databases

Image databases including *annotations* (typically some kind of high level information) are essential to the development of *supervised* machine learning methods for image analysis.

Annotations

- Image class
- Measure(s) obtained from the image
- Position of objects within the image
- Segmentation

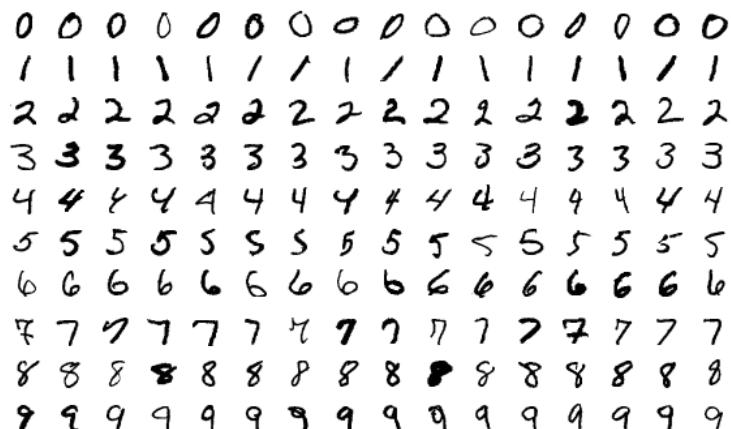
6 / 59

MNIST database [Lecun et al., 1998]

- The Modified National Institute of Standards and Technology (MNIST) database contains 60 000 training images of hand-written digits, and 10,000 test images.
- Image size: 28×28
- It has been used since 1998
- Human performance on a similar database (NIST) is reported to be around 1.5% error [Simard et al., 1993]
- Best methods, based on convolutional neural networks, give around 0.21% test error.

7 / 59

MNIST database



Credits: Images from MNIST assembled by Josef Stepan (licensed under CC BY-SA 4.0)

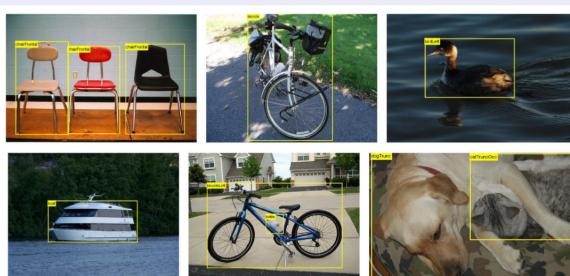
8 / 59

Pascal VOC project [Everingham et al., 2010, Everingham et al., 2014]

This project organized a challenge from 2005 to 2012, divided into several tasks, including an image classification task.

Pascal VOC image classification task (2012)

Train/val: 11 540 images where the presence of 20 categories of objects was annotated. The test dataset is unknown and tests are run online (still available).



Credits: From [Everingham et al., 2014]

9 / 59

ImageNet project

Since 2010, ImageNet organizes an annual challenge: The ImageNet Large Scale Visual Recognition Challenge (ILSVRC), that constitutes a breakthrough in the design of image analysis challenges by its size.

Image classification task (since 2012)

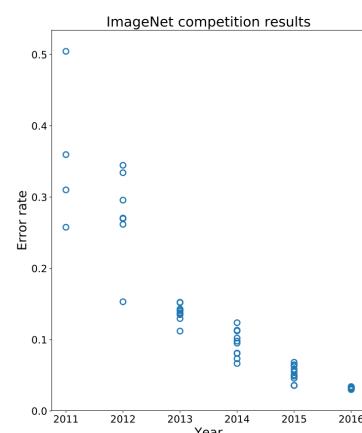
- Training: 1 281 167; validation: 50 000; test: 100 000.
- 1 000 classes (90 dog breeds!).

10 / 59

ImageNet project



Examples from the *acoustic guitar* class

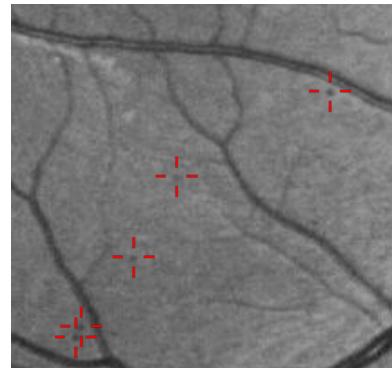


Credits: Wikipedia (CC BY-SA 4.0)

11 / 59

Image processing approach

- Build a geometrical model for the objects of interest
- Implement this model using image processing operators



Detail of eye fundus image with microaneurysms to be detected

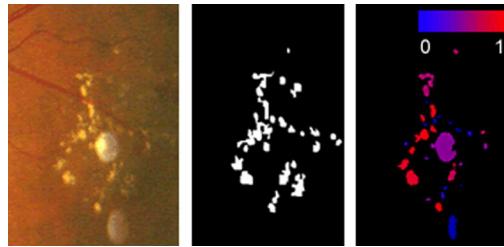
- + This approach works correctly when the objects are not too complex.
- If objects are difficult to model, machine learning methods can bring a solution.

Credits: [zhang et al., 2011]

12 / 59

Classical machine learning approach

- Compute features from the image
- Apply machine learning to those features



Exudates segmentation: original image, ground-truth and candidates with associated probabilities obtained with machine learning

- + Works well with the right features
 - An expert is required to define those features - and this can be a long process
 - Annotated data is required

Credits: [Zhang et al., 2014]

13 / 59

Deep learning approach

- Directly take as input the image pixels
 - The network is supposed to build its own features
- + Good (impressive!) results
 - A large amount of annotated data is required

14 / 59

Some accomplishments

ImageNet Large Scale Visual Recognition Challenge (ILSVRC)

2012: *AlexNet* [Krizhevsky et al., 2012] won this challenge by a large margin

The database contains more than 1 million training images, belonging to 1000 different classes.

15 / 59

Some accomplishments (cont.)

- 2011: first super-human performance, IJCNN 2011 traffic sign recognition contest [Cireşan et al., 2011]
- 2012: visual object detection (Mitosis detection in breast cancer histology [Cireşan et al., 2013])
- 2012: segmentation competition (neuronal membranes in electron microscopy images [Ciresan et al., 2012])
- 2016: AlphaGo beats Lee Sedol, one of the best go players, in a 5-game match

16 / 59

Deep learning image applications

- classification
- object localization
- semantic segmentation
- instance segmentation
- transformation (filtering, in-painting, editing, colorization...)
- quantification
- compression
- image caption generation
- 2D to 3D (stereo matching, 3D reconstruction, ...)
- motion estimation
- Style transfer
- Anomalous image detection
- Image generation

17 / 59

Convolutional neural networks in deep learning

- They are pivotal to many of the successes achieved by neural networks these recent years
- They are interesting for dealing with regular structured data, such as images (or board games!)

Acronyms

CNN and *ConvNet*

Essential milestones

- 1979: Neocognitron (CNN architecture)
[Fukushima, 1979, Fukushima, 1980]
- 1989: Backpropagation applied to CNNs [LeCun et al., 1989]
- 2006, 2010: GPU implementation
[Chellapilla et al., 2006, Cireşan et al., 2010]
- 2010: Availability of large databases (ImageNet, ...)

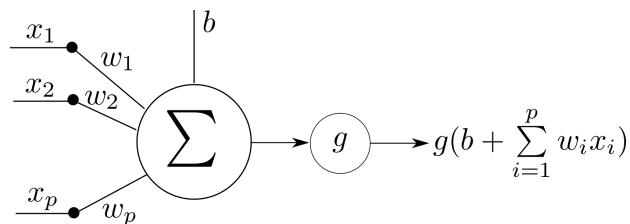
18 / 59

Contents

- 1 Introduction
- 2 Application of fully-connected networks to image classification
- 3 From fully-connected layers to convolutional layers
- 4 Building convolutional networks
- 5 Conclusion

19 / 59

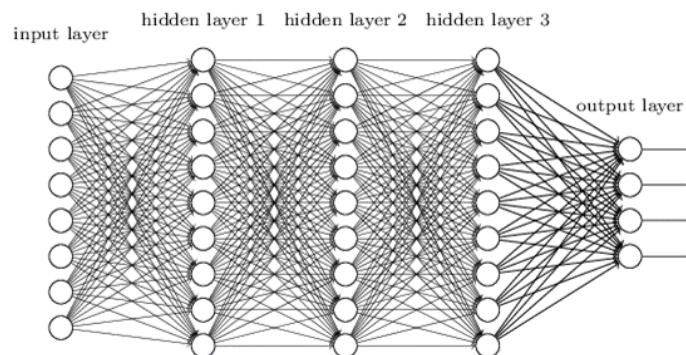
Reminder: Artificial neuron



- b, w_1, \dots, w_n are the neuron parameters, to be learnt
- g is the activation or transfer function

20 / 59

Reminder: Neural network

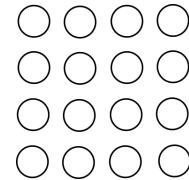
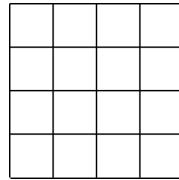


(from <http://www.jtoy.net>)

21 / 59

Input image, input neurons

In the scalar case (single-valued images), each input pixel is considered as an input neuron.



22 / 59

Image classification problem

Classification problem:

- Input: image \mathbf{x}
- Output: class $y \in \{label_1, label_2, \dots, label_q\}$

Class coding

Often, classes are denoted by integers, but this is only a coding commodity. For instance, it would be meaningless to use a regression approach for this problem.

23 / 59

Class coding

If there are q possible classes, then a class will be coded as a vector \mathbf{y} of length q . If its class is r then for $0 \leq i < q$:

$$\mathbf{y}[i] = \begin{cases} 1, & \text{if } i = r \\ 0, & \text{otherwise} \end{cases}$$

Example with 4 classes

- Label 0 $\mapsto [1, 0, 0, 0]$
- Label 1 $\mapsto [0, 1, 0, 0]$
- Label 2 $\mapsto [0, 0, 1, 0]$
- Label 3 $\mapsto [0, 0, 0, 1]$

24 / 59

Image classification with a neural network

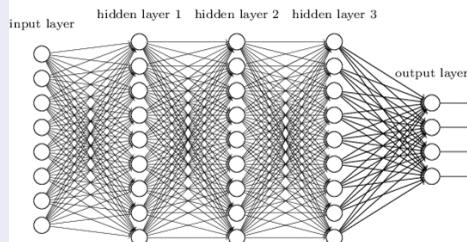
Input

The input image, containing p pixels, is transformed into a vector of length p .

Output

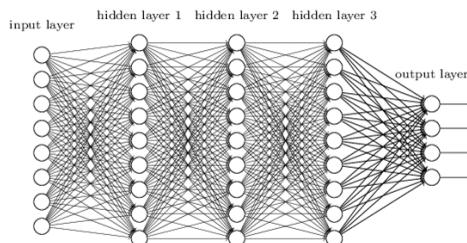
For q classes, the output will be a vector of length q .

Example: image of size 4×2 , 4 possible classes



25 / 59

Image classification using fully-connected layers



- A small image contains at least 100 000 pixels
- The number of parameters between two layers of that size is $10^5 \times (10^5 + 1)!$
- This approach is not feasible...
- Moreover, this approach does not take into account the local structure of images.

26 / 59

Activations

Different activations (typically ReLU) can be used in the intermediate layers.

Concerning the last layer: Given that the aim is a vector containing zeros except for a one, two designs are commonly used:

- Use a sigmoid as last activation
- Use any activation, but follow it by a soft-max operator

27 / 59

Softmax operator

Definition

The softmax operator $\sigma : \mathbb{R}^d \rightarrow \mathbb{R}^d$ is given by:

$$\forall \mathbf{x} \in \mathbb{R}^d, \forall k \in \{1, \dots, d\} : \quad \sigma(\mathbf{x})_k = \frac{e^{\mathbf{x}_k}}{\sum_{i=1}^d e^{\mathbf{x}_i}}$$

Some properties

- $0 < \sigma(\mathbf{x})_k < 1$
- $\sum_{i=1}^d \sigma(\mathbf{x})_i = 1$

Example

$$\mathbf{x} = \begin{pmatrix} 10.1 \\ 0 \\ -4.3 \\ 1.33 \end{pmatrix} \quad \sigma(\mathbf{x}) \approx \begin{pmatrix} 0.9998 \\ 0.000041 \\ 0.00000056 \\ 0.00016 \end{pmatrix}$$

28 / 59

Loss function for classification: cross-entropy

The preferred loss function for classification is cross-entropy:

For \mathbf{y} in $[0, 1]^d$ and $\hat{\mathbf{y}}$ in $]0, 1[^d$:

$$H(\mathbf{y}, \hat{\mathbf{y}}) = \sum_{i=1}^d \mathbf{y}_i \log(\hat{\mathbf{y}}_i)$$

For this to work, each $\hat{\mathbf{y}}_i$ must be strictly positive. Therefore when used in a NN the cross-entropy has to be preceded by a convenient operator, such as a sigmoid or a softmax.

Note that the binary cross-entropy we saw during the first lecture is a particular case of cross-entropy.

29 / 59

Conclusion on fully-connected networks for image classification

Fully connected layers:

- scale badly to large size images
- do not take into account the local structure of images

Today:

- Fully-connected networks are almost never used for image analysis.
- Fully-connected layers are only used in the middle (auto-encoders) or at the end (classification) of the pipeline.

30 / 59

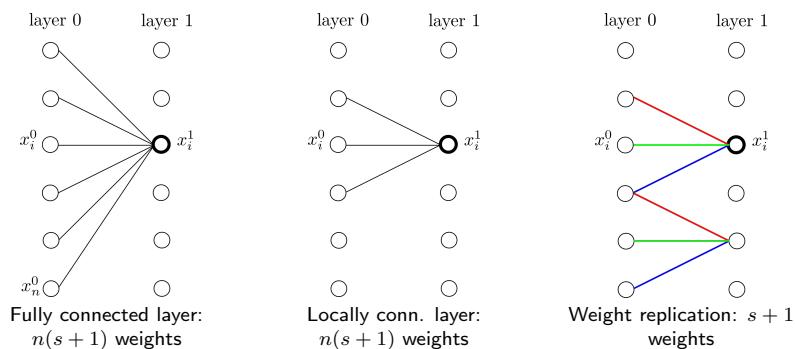
Contents

- 1 Introduction
- 2 Application of fully-connected networks to image classification
- 3 From fully-connected layers to convolutional layers
- 4 Building convolutional networks
- 5 Conclusion

31 / 59

Towards convolutional layers

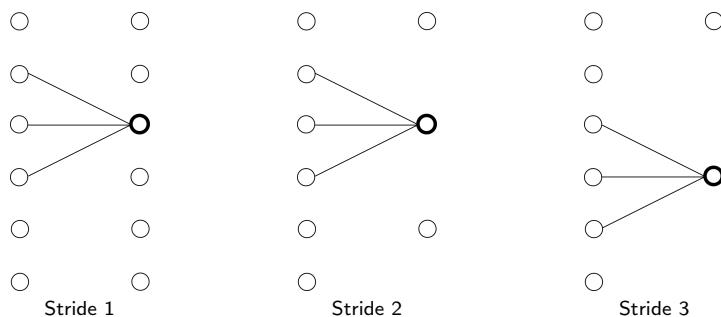
For illustration purposes, in the following slides images and filters will be displayed as rows of neurons - these can be seen as 1D arrays or as sections of 2D arrays.



32 / 59

Stride

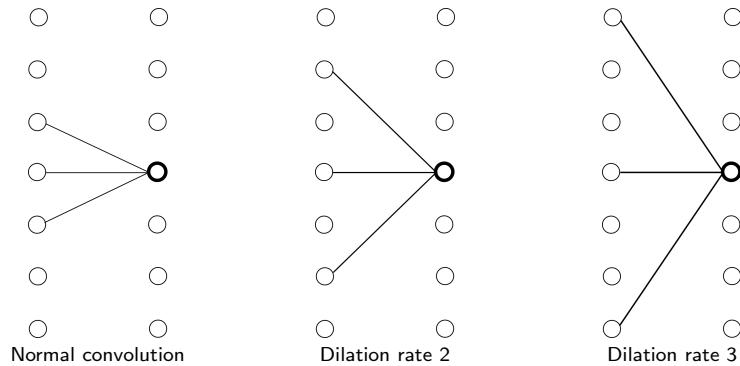
A convolutional layer can at the same time downsample the image by applying a sampling step, or *stride*.



33 / 59

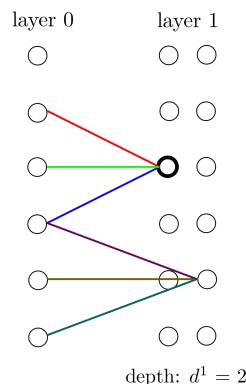
Dilated convolutions

Dilated convolutions are used to increase the size of the receptive field of the network.



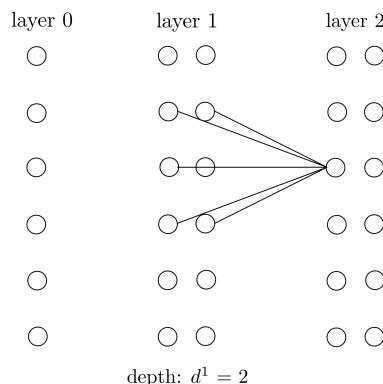
34 / 59

Several filters in the same convolutional layer



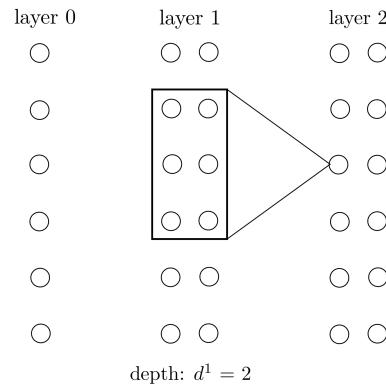
35 / 59

Several filters in the same convolutional layer



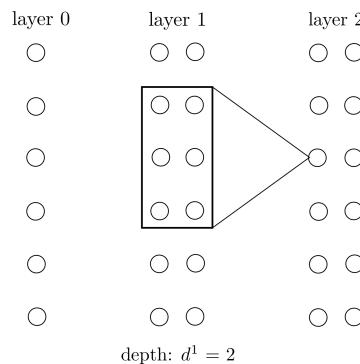
36 / 59

Several filters in the same convolutional layer



37 / 59

Consequences on the parameter number

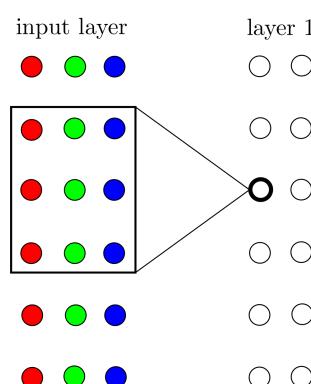


- How many parameters do we have in layer 1?
- $d^1 \times (s + 1)$
- In layer 2?
- $d^2 \times (d^1 \times s + 1)$

38 / 59

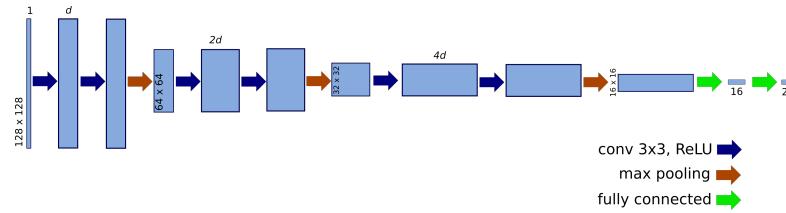
Multi-valued images

An input image with p channels (for instance a colour image with 3 channels) can be represented by an input layer of depth 3



39 / 59

1D representations

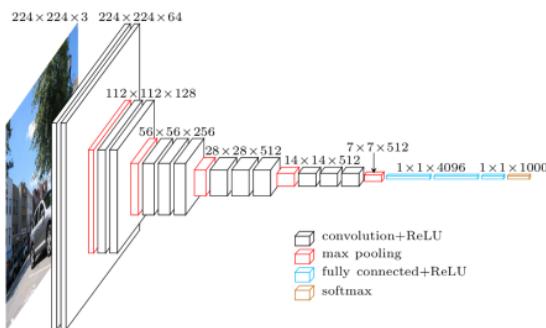


This NN was used to estimate the position of the center of the macula on fundus images.

Credits: NN is work of Robin Alais et al.
Fundus image by Mikael Häggström, used with permission (CC0).

40 / 59

2D representations



This network is used for image classification tasks.

Credits: VGG16 (From <https://www.cs.toronto.edu/~frossard/post/>)

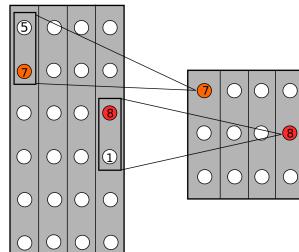
41 / 59

Contents

- 1 Introduction
- 2 Application of fully-connected networks to image classification
- 3 From fully-connected layers to convolutional layers
- 4 Building convolutional networks
- 5 Conclusion

Max pooling

- Convolutional networks often contain subsampling steps. A common way of doing this today is by using *max pooling* layers with stride 2.
- Sampling is only applied along the spatial dimensions, not along the dimension of the filters.

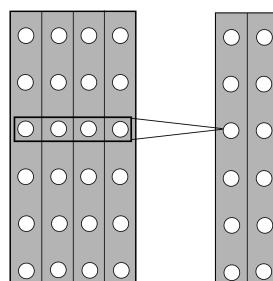


Note however a current trend that consists in using convolutional layers with a stride of 2

43 / 59

Dimension reduction

1×1 convolutions are used to reduce the number of filters - this is called by some authors *dimension reduction*.



44 / 59

Main components of a convolutional neural network

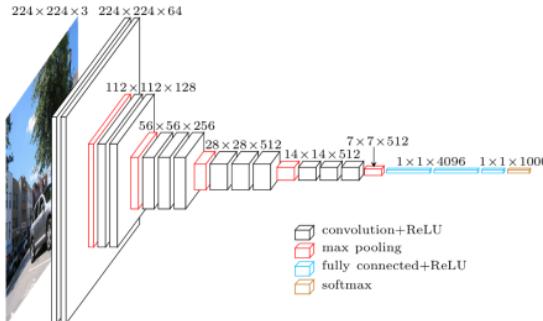
Many successful architectures, especially for image classification, follow the same pattern:

- ① Several iterations of: One or several convolutional layers, with increasing depth, followed by max pooling
- ② A few fully connected layers

45 / 59

VGGnet

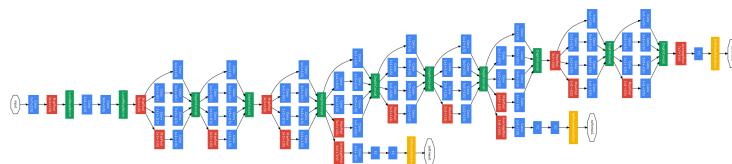
- Proposed by K. Simonyan and A. Zisserman from the University of Oxford [Simonyan and Zisserman, 2014]
- Runner-up in the ImageNet Large Scale Visual Recognition Competition (ILSVRC) in 2014.
- Number of parameters (VGG16): 138 million.



Credits: VGG16 (From
<https://www.cs.toronto.edu/~frossard/post/>)
46 / 59

GoogLeNet (a.k.a. Inception v1)

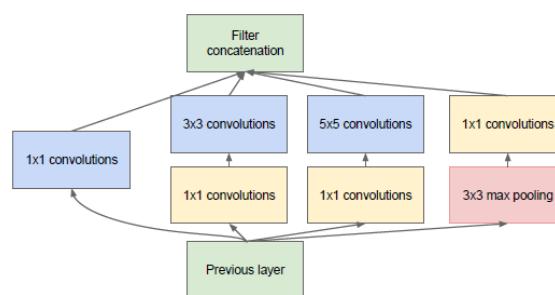
- Winner of the ImageNet Large Scale Visual Recognition Competition (ILSVRC) in 2014.
- Number of parameters: *only* 5 million.



ResNet won the following year...

Credits: From [Szegedy et al., 2014]
47 / 59

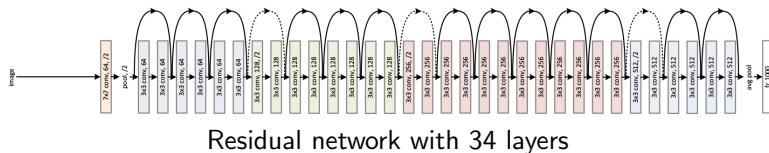
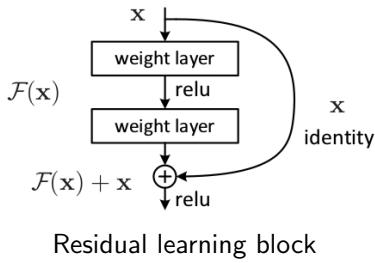
Inception module



Credits: From [Szegedy et al., 2014]
48 / 59

ResNet

- Winner of the ImageNet Large Scale Visual Recognition Competition (ILSVRC) in 2015.
 - The authors tested up to 1202 layers. They reported no training difficulties, but overfitting [He et al., 2015]



Credits: From [He et al., 2015]

49 / 59

Current trends

- Small convolutions (3×3)
 - Dimension reduction using 1×1 convolutions
 - Increasing number of layers
 - Skip connections

VGG, GoogLeNet and ResNet (and their variants) are still among the most used architectures for image classification and other related tasks.

50 / 59

Some deep learning libraries

Deep learning is a very competitive domain, where code sharing is very common.

- Tensorflow, by Google (Apache licence)
 - PyTorch, Torch (Facebook - BSD licence)
 - Caffe (Univ. of California, Berkeley - BSD licence)
 - Microsoft Cognitive Toolkit (MIT licence)
 - MatConvNet (for MatLab users)
 - Theano (Montreal Institute for Learning Algorithms; not maintained anymore)

Comments

- Most of these libraries are distributed with very permissive licences
 - Most of them use Python as prototyping language
 - Keras is a very easy to use interface to Tensorflow, Theano and CNTK

51 / 59

Contents

- 1 Introduction
- 2 Application of fully-connected networks to image classification
- 3 From fully-connected layers to convolutional layers
- 4 Building convolutional networks
- 5 Conclusion

52 / 59

A revolution in image analysis

- Deep learning has brought an undeniable break-through in image analysis (as in other fields)
- A significant part of research efforts in image analysis today is based on deep learning
- Its applications are ubiquitous

53 / 59

Limitations

For a deep-learning solution to work, you need:

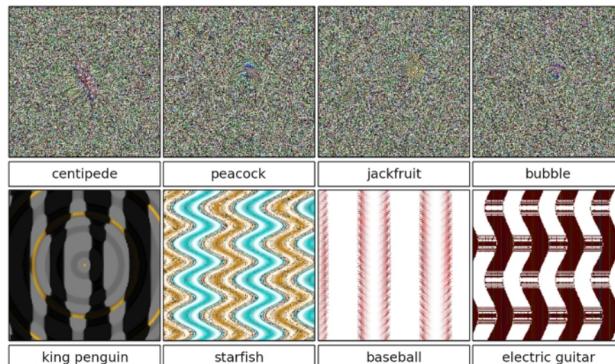
- Enough annotated data
- A lot of fiddling (different architectures; hyper-parameters; optimization)
- One (or, even better, several) powerful GPUs

Moreover, these models lack interpretability.

54 / 59

ConvNets can be fooled

Deep learning can produce astonishing results
[Nguyen et al., 2015]...



55 / 59

References I

- [Chellapilla et al., 2006] Chellapilla, K., Puri, S., and Simard, P. (2006). High Performance Convolutional Neural Networks for Document Processing. In *Tenth International Workshop on Frontiers in Handwriting Recognition*. Suvisoft.
- [Ciresan et al., 2012] Ciresan, D., Giusti, A., Gambardella, L. M., and Schmidhuber, J. (2012). Deep Neural Networks Segment Neuronal Membranes in Electron Microscopy Images. In Pereira, F., Burges, C. J. C., Bottou, L., and Weinberger, K. Q., editors, *Advances in Neural Information Processing Systems 25*, pages 2843–2851. Curran Associates, Inc.
- [Cireşan et al., 2011] Cireşan, D., Meier, U., Masci, J., and Schmidhuber, J. (2011). A committee of neural networks for traffic sign classification. In *Neural Networks (IJCNN), The 2011 International Joint Conference on*, pages 1918–1921. IEEE.
- [Cireşan et al., 2013] Cireşan, D. C., Giusti, A., Gambardella, L. M., and Schmidhuber, J. (2013). Mitosis Detection in Breast Cancer Histology Images with Deep Neural Networks. In *Medical Image Computing and Computer-Assisted Intervention – MICCAI 2013*, pages 411–418. Springer, Berlin, Heidelberg.
- [Cireşan et al., 2010] Cireşan, D. C., Meier, U., Gambardella, L. M., and Schmidhuber, J. (2010). Deep, Big, Simple Neural Nets for Handwritten Digit Recognition. *Neural Computation*, 22(12):3207–3220.

56 / 59

References II

- [Everingham et al., 2014] Everingham, M., Eslami, S. M. A., Gool, L. V., Williams, C. K. I., Winn, J., and Zisserman, A. (2014). The Pascal Visual Object Classes Challenge: A Retrospective. *International Journal of Computer Vision*, 111(1):98–136.
- [Everingham et al., 2010] Everingham, M., Van Gool, L., Williams, C. K. I., Winn, J., and Zisserman, A. (2010). The Pascal Visual Object Classes (VOC) Challenge. *International Journal of Computer Vision*, 88(2):303–338.
- [Fukushima, 1979] Fukushima, K. (1979). Neural Network Model for a Mechanism of Pattern Recognition Unaffected by Shift in Position- Neocognitron. *ELECTRON. & COMMUN. JAPAN*, 62(10):11–18.
- [Fukushima, 1980] Fukushima, K. (1980). Neocognitron: A self-organizing neural network model for a mechanism of pattern recognition unaffected by shift in position. *Biological Cybernetics*, 36(4):193–202.
- [He et al., 2015] He, K., Zhang, X., Ren, S., and Sun, J. (2015). Deep Residual Learning for Image Recognition. *arXiv:1512.03385 [cs]*. arXiv: 1512.03385.
- [Krizhevsky et al., 2012] Krizhevsky, A., Sutskever, I., and Hinton, G. E. (2012). ImageNet Classification with Deep Convolutional Neural Networks. In Pereira, F., Burges, C. J. C., Bottou, L., and Weinberger, K. Q., editors, *Advances in Neural Information Processing Systems 25*, pages 1097–1105. Curran Associates, Inc.

57 / 59

References III

- [LeCun et al., 1989] LeCun, Y., Boser, B., Denker, J. S., Henderson, D., Howard, R. E., Hubbard, W., and Jackel, L. D. (1989). Backpropagation Applied to Handwritten Zip Code Recognition. *Neural Computation*, 1(4):541–551.
- [Lecun et al., 1998] Lecun, Y., Bottou, L., Bengio, Y., and Haffner, P. (1998). Gradient-based learning applied to document recognition. *Proceedings of the IEEE*, 86(11):2278–2324.
- [Nguyen et al., 2015] Nguyen, A., Yosinski, J., and Clune, J. (2015). Deep neural networks are easily fooled: High confidence predictions for unrecognizable images. In *2015 IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 427–436.
- [Simard et al., 1993] Simard, P., LeCun, Y., and Denker, J. S. (1993). Efficient pattern recognition using a new transformation distance. In *Advances in neural information processing systems*, pages 50–58.
- [Simonyan and Zisserman, 2014] Simonyan, K. and Zisserman, A. (2014). Very deep convolutional networks for large-scale image recognition. *arXiv preprint arXiv:1409.1556*.
- [Szegedy et al., 2014] Szegedy, C., Liu, W., Jia, Y., Sermanet, P., Reed, S., Anguelov, D., Erhan, D., Vanhoucke, V., and Rabinovich, A. (2014). Going Deeper with Convolutions. *arXiv:1409.4842 [cs]*. arXiv: 1409.4842.

58 / 59

References IV

- [zhang et al., 2011] zhang, x., thibault, g., and decenciere, e. (2011). Application of the Morphological Ultimate Opening to the Detection of Microaneurysms on Eye Fundus Images from Clinical Databases. In *ICS'13*.
- [Zhang et al., 2014] Zhang, X., Thibault, G., Decencière, E., Marcotegui, B., Laÿ, B., Danno, R., Cazuguel, G., Quellec, G., Lamard, M., Massin, P., Chabouis, A., Victor, Z., and Erginay, A. (2014). Exudate detection in color retinal images for mass screening of diabetic retinopathy. *Medical Image Analysis*, 18(7):1026–1043.

59 / 59